Qwen3.8-27B native DFlash2 · 7 draft tokens · 64 output tokens · chunk 128

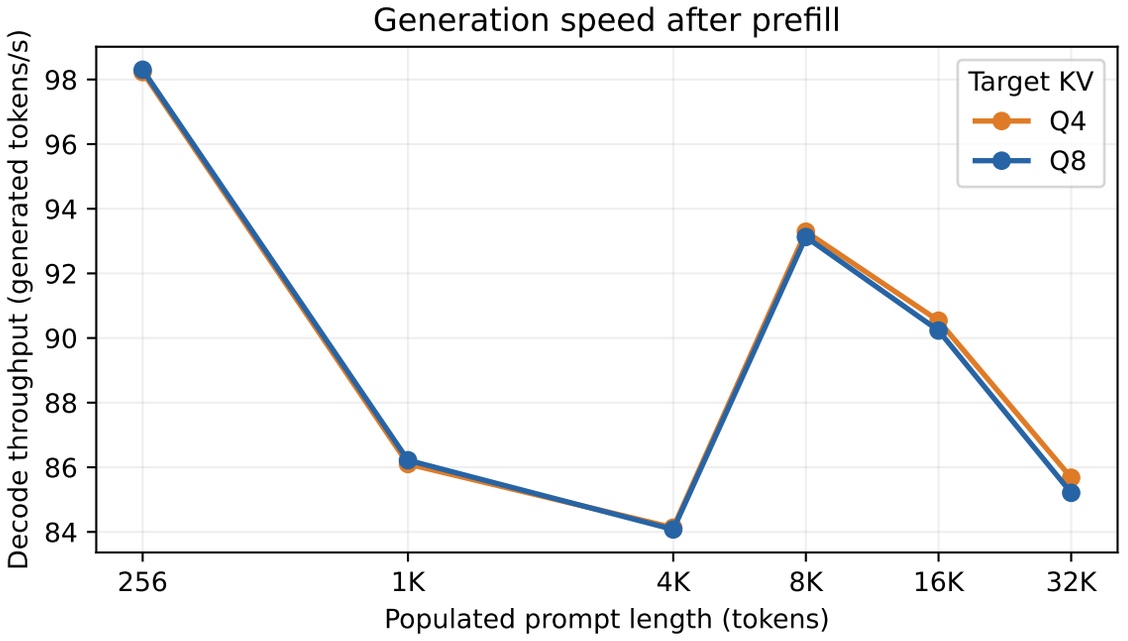

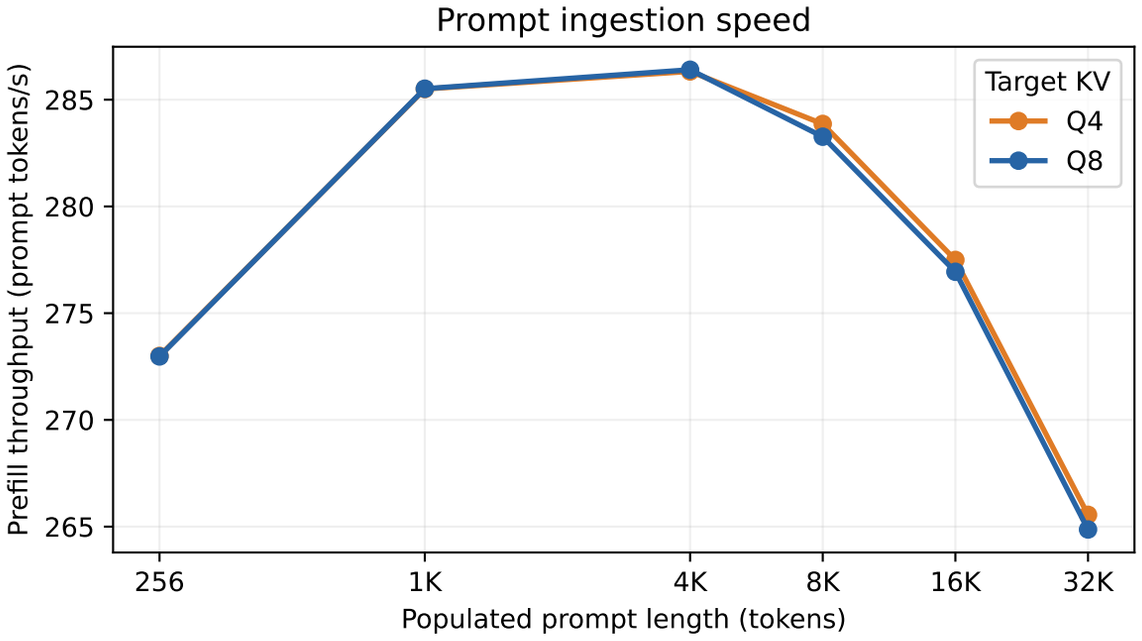

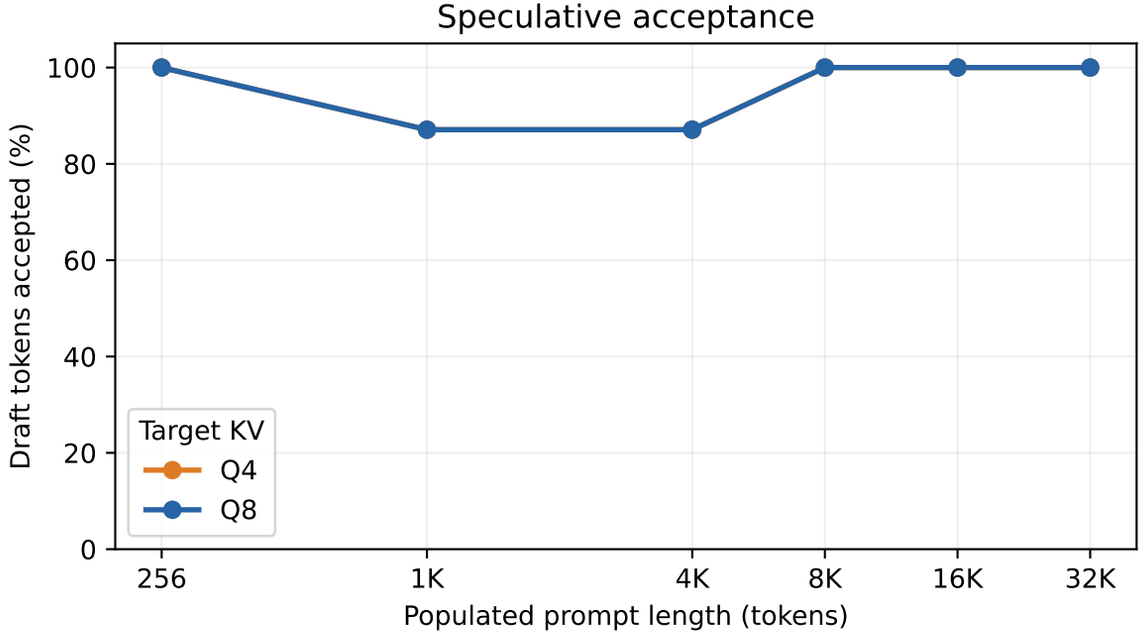

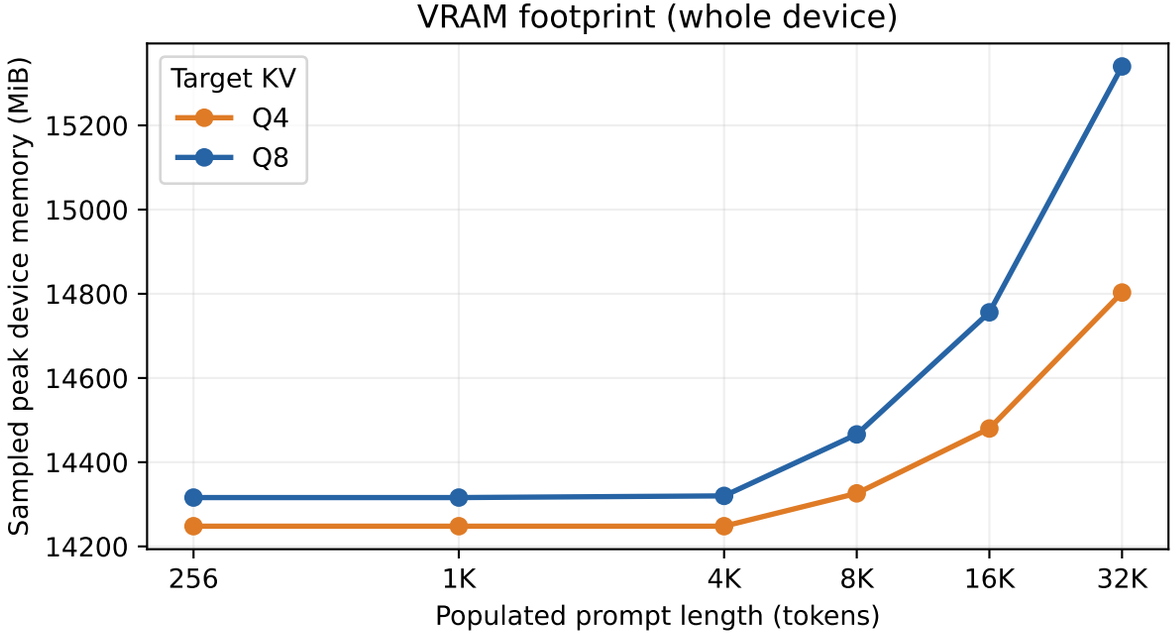

Fresh process per point. Prompts repeat the selected fixture body; this can affect draft acceptance. One measured repetition per point; variability is not estimated. VRAM is sampled whole-device nvidia-smi usage.